

Dependency Paging

Tal Zohar

May 4, 2026

1 The Dependency Paging Problem

There are n pages and a cache of size k . Additionally, we are provided with a dependency DAG G . Let $G[p]$ be the set of dependencies of page p . A valid solution must ensure that if page p is in cache, then all pages in $G[p]$ are also in the cache. Let l denote the width of G - the maximum antichain length.

We analyze the problem assuming pages are requested one by one, such that every request arrives only after all of its children (dependencies) have been requested. This topological request order simplifies the algorithm description and proofs without loss of generality.

2 Marking Algorithms and Phase Structure

To analyze both the deterministic and randomized algorithms, we utilize the classical framework of *marking algorithms*, adapted to respect dependency constraints.

A marking algorithm processes a request sequence in a series of disjoint phases. At the beginning of a phase, all pages are considered “unmarked.” As requests arrive, the requested pages (and their dependencies) are “marked.” The fundamental rule of a marking algorithm is that it **never evicts a marked page**. When the cache is full and an eviction is required, the algorithm is forced to choose an unmarked page. A phase naturally terminates when an eviction is needed but the cache is entirely full of marked pages. At this point, all marks are cleared, and a new phase begins.

Consequently, the sequence is divided into phases containing exactly k distinct pages, where k represents the size of the cache. Let m denote the number of new pages requested in a phase (pages that were not in the cache at the start of the phase).

Lemma 1 (OPT Lower Bound). *The optimal offline algorithm (OPT) incurs an amortized cost of at least $m/2$ misses per phase.*

Proof. Because OPT has a cache capacity of exactly k , and each phase requests exactly k distinct pages including m entirely new pages, OPT must incur at least m misses across the current and previous phases combined. Thus, its amortized cost per phase is at least $m/2$. $\square$

By Lemma 1, any algorithm that incurs $O(m \cdot f(l))$ misses per phase is $O(f(l))$ -competitive.

We partition the DAG G into l disjoint strongly-ordered chains $C_1, \dots, C_l$, which is possible by Dilworth’s theorem.

Definition: A page is considered **marked** if it has been requested during the current phase.

Definition: A chain is considered **marked** if there is no unmarked page of the chain currently residing in the cache (i.e., all chain pages in the cache are marked).

Definition: A page in the cache is an **eviction candidate** if it is a maximal unmarked cached page (meaning it is unmarked, and no other page currently in the cache is dependent on it). Because requests follow a topological order, these are the only pages that can be safely evicted without violating dependency constraints.

Observation 1. Equivalence of Phase Boundaries: *By definition, a phase ends when all k pages in the cache are marked. Because the cache capacity is exactly k , this is the exact moment when no unmarked pages remain in the cache, and consequently, all l chains become marked. Therefore, a phase runs exactly as long as there are valid eviction candidates in the cache.*

Observation 2. *Within each chain, the marked pages reside at the bottom, while the evicted pages reside at the top.*

3 A Deterministic $O(\min(k, l))$ -Competitive Algorithm

Standard LRU preserves dependency constraints and is $O(k)$ -competitive, as shown in [1]. By a tighter analysis, we can show that a generic dependency-aware marking algorithm is $O(\min(k, l))$ -competitive.

Claim 1. *Standard LRU is a valid implementation of the generic deterministic dependency-aware marking algorithm.*

Algorithm 1 describes this generic deterministic strategy. Whenever the cache is full, it simply evicts *any* valid eviction candidate.

Algorithm 1 Deterministic Marking Algorithm for Dependency Paging

```

1: Initialize: Cache  $S \leftarrow \emptyset$ , Marked set  $M \leftarrow \emptyset$ 
2: for each requested page  $p$  in the sequence do
3:   if  $|M \cup \{p\}| > k$  then                                 $\triangleright$  Phase ends when we attempt to mark beyond capacity
4:     End current phase, start new phase
5:      $M \leftarrow \emptyset$                                         $\triangleright$  Unmark all pages
6:     Mark all pages in  $S$  that are needed for  $p$ 
7:   end if
8:    $M \leftarrow M \cup \{p\} \cup G[p]$                            $\triangleright$  Mark the requested page and its dependencies
9:   while  $|S \cup \{p\}| > k$  do                                 $\triangleright$  Cache is full, eviction required
10:    Find an eviction candidate  $q \in S$ 
11:     $S \leftarrow S \setminus \{q\}$                                  $\triangleright$  Evict page  $q$ 
12:  end while
13:  if  $p \notin S$  then
14:     $S \leftarrow S \cup \{p\}$                                      $\triangleright$  Fetch the requested page into the cache
15:  end if
16: end for

```

To prove the competitiveness, we bound the number of misses on “old pages” (pages that were in the cache at the start of the phase).

Lemma 2 (Bound on Displaced Pages). *At any given moment during a phase, the total number of old pages that are currently evicted from the cache is bounded by m .*

Proof. Because the cache maintains a fixed capacity of k , every currently evicted old page must have been perfectly replaced by a “new page” requested during this phase. Since at most m new pages are requested in the entire phase, the cache can physically hold at most m new pages at any given time. Thus, the total number of evicted old pages is bounded by m . $\square$

Lemma 3 (Chain Marking Upon Miss). *Let p be an old page belonging to chain C_j that was evicted during the current phase. At the exact moment p is subsequently requested and brought back into the cache, C_j becomes a marked chain.*

Proof. Because C_j is a totally ordered chain, any other page $y \in C_j$ must either be a dependency of p , or y must depend on p .

- If y is a dependency of p , the algorithm guarantees that when p is marked, all of its dependencies currently in the cache are also marked. Thus, y is marked.
- If y depends on p , then y requires p to be in the cache. When p was originally evicted, it was an eviction candidate, meaning no cached page depended on it. Since p has remained evicted until now, no page depending on p (including y) could have been requested or fetched into the cache, as doing so would require fetching p first. Thus, y is currently not in the cache.

Therefore, every page of C_j currently in the cache is marked, making C_j a marked chain. $\square$

Theorem 1. *The generic deterministic marking algorithm is $O(\min(k, l))$ -competitive.*

Proof. We are interested in bounding the total number of misses on old pages. An old page incurs a miss if and only if it is evicted and subsequently refetched.

First, observe that once an old page is refetched, it becomes marked. By the fundamental rule of marking algorithms, marked pages are never evicted. Thus, each of the at most k old pages originally in the cache can incur a maximum of one miss per phase. This bounds the total number of misses on old pages strictly by k .

Second, we can bound the misses using the chain structure. By Lemma 3, the very first time an evicted old page from C_j is refetched, C_j becomes a marked chain. A marked chain never becomes unmarked again for the remainder of the phase.

Because the algorithm only evicts unmarked pages (eviction candidates), no page from C_j can be evicted once C_j becomes marked. Therefore, all evictions from C_j must occur *before* C_j suffers its first miss.

By Lemma 2, the maximum number of pages from C_j that can be simultaneously evicted at any given moment is bounded by the total number of evicted old pages in the entire system, which is m . Since C_j cannot be evicted from after its first miss, the maximum number of misses C_j can incur is exactly the number of its pages that were evicted prior to this moment, which is at most m .

Summing this bound over all l chains, the total number of misses on old pages is also bounded by $m \cdot l$.

Combining these two independent bounds, the total number of misses on old pages is at most $\min(k, m \cdot l) \leq m \cdot \min(k, l)$ (since $m \geq 1$).

The algorithm also incurs exactly m misses to fetch the m new pages. Thus, the total cost for the phase is at most $m + m \cdot \min(k, l) = m(1 + \min(k, l))$. By Lemma 1, OPT incurs an amortized cost of at least $m/2$ per phase. Therefore, the generic deterministic marking algorithm achieves an overall competitive ratio of $O(\min(k, l))$. $\square$

4 A Randomized $O(\min(\log k, \log l))$ -Competitive Algorithm

The paper [1] provides an expected $O(\min(\log k, \log l))$ -competitive randomized algorithm. Here, we provide an alternative, simpler analysis to establish this exact $O(\min(\log k, \log l))$ competitive ratio directly.

We present a randomized algorithm (Algorithm 2) that restricts how the eviction candidates are chosen. By ensuring evictions are distributed symmetrically across the chains, we prevent the adversary from stacking m misses on a single chain.

4.1 Blocked Chains and Algorithm Description

We reuse the chains $C_1, \dots, C_l$ and the definition of an eviction candidate. To formalize the randomized selection, we define how a chain specifies its desired eviction target.

Definition: For any unmarked chain, its **next eviction candidate** is defined as its highest unmarked page currently residing in the cache. Note that a chain's next eviction candidate might not be a global *eviction candidate*, as there could be cached pages *outside* the chain that depend on it.

We introduce the concept of a “blocked chain” to systematically determine how actual eviction candidates are chosen.

Definition: An unmarked chain is considered **blocked** on a page p if the following two conditions are met:

1. p is an eviction candidate.
2. p is an ancestor of (or equal to) the chain's next eviction candidate.

Claim 2. *Every unmarked chain must be blocked by some page (can be multiple).*

Proof. Let $\tilde{C}$ be an unmarked chain. By definition, it has a next eviction candidate $x \in \tilde{C}$ currently residing in the cache. If we trace the dependency graph upwards from x to its ancestors currently in the cache, we will eventually reach a maximal unmarked cached page q (one with no cached dependents). This page q is an eviction candidate. Because q is either equal to x or an ancestor of x , the chain $\tilde{C}$ is blocked on q . $\square$

Eviction Rule: When we need to evict a page, we draw an unmarked chain $\tilde{C}$ uniformly at random. We select some page q such that $\tilde{C}$ is blocked on q (by Claim 2, we are guaranteed such a page exists), and we evict q .

Algorithm 2 Randomized Marking Algorithm for Dependency Paging

```
1: Initialize: Cache  $S \leftarrow \emptyset$ , Marked set  $M \leftarrow \emptyset$ 
2: Partition  $G$  into  $l$  disjoint chains  $C_1, \dots, C_l$  ▷ via Dilworth's Theorem
3: for each requested page  $p$  in the sequence do
4:   if  $|M \cup \{p\}| > k$  then ▷ Phase ends when we attempt to mark beyond capacity
5:     End current phase, start new phase
6:      $M \leftarrow \emptyset$  ▷ Unmark all pages
7:     Mark all pages in  $S$  that are needed for  $p$ 
8:   end if
9:    $M \leftarrow M \cup \{p\} \cup G[p]$  ▷ Mark the requested page and its dependencies
10:  while  $|S \cup \{p\}| > k$  do ▷ Cache is full, eviction required
11:    Draw an unmarked chain  $\tilde{C}$  uniformly at random
12:    Select a page  $q \in S$  such that  $\tilde{C}$  is blocked on  $q$ 
13:    ▷ ( $q$  is an eviction candidate, and ancestor of  $\tilde{C}$ 's next eviction candidate)
14:     $S \leftarrow S \setminus \{q\}$  ▷ Evict page  $q$ 
15:  end while
16:  if  $p \notin S$  then
17:     $S \leftarrow S \cup \{p\}$  ▷ Fetch the requested page into the cache
18:  end if
19: end for
```

4.2 Conceptual Accounting Scheme and Analysis

To bound the expected number of cache misses, we introduce a conceptual accounting scheme using “coins”. It is important to note that these coins are strictly an analytical tool used for the proof; the algorithm does not track or utilize them during its execution.

Coin Distribution Rules:

- **Giving Coins:** Whenever the algorithm evicts a page q by randomly selecting an unmarked chain $\tilde{C}$, we conceptually give one coin to chain $\tilde{C}$. We define the evicted page q to be the page “in charge” of this coin.
- **Removing Coins:** If a page q that is in charge of a coin is subsequently requested (and thus marked), we remove its associated coin from its chain. This scenario occurs exclusively when an evicted old page is refetched into the cache.

Based on this accounting scheme and the structural properties established in Section 2, we establish the following claims:

Claim 3 (Bound on Total Coins). *At any given moment, the total number of coins in the system is bounded by m .*

Proof. By our accounting rules, a coin corresponds exactly to an old page that is currently evicted. By Lemma 2, the number of currently evicted old pages is bounded by m . Thus, the total number of coins is bounded by m . $\square$

Claim 4 (Coin Removal). *We only remove coins from marked chains.*

Proof. A coin is removed from a chain C' exactly when the page p in charge of it is requested. We must show that at this moment, C' is a marked chain (i.e., it has no unmarked pages in the cache).

When p was originally evicted and took charge of the coin, C' was blocked on p . By definition, this means C' had a next eviction candidate x (its highest unmarked cached page at that time), and p was an ancestor of x (or equal to x).

Fast forward to the moment p is requested and the coin is removed. Because p is requested, p is immediately marked. The algorithm guarantees that when a page is marked, all its dependencies currently in the cache are also marked. Because p is an ancestor of x , x is a dependency of p . Thus, x becomes marked.

Could there be some other page $y \in C'$ in the cache that is unmarked? Because C' is a totally ordered chain of dependencies, y must be either a dependency of x , or x must be a dependency of y .

- If y is a dependency of x , then by transitivity, y is a dependency of p . For the same reason x is marked, y must also be marked.
- If x is a dependency of y , then y is an ancestor of x . However, recall that when p was evicted, it was an eviction candidate, meaning it had *no* cached dependents (no cached pages depending on it). Because C' is a totally ordered chain, y and p must be ordered. Since both y and p are ancestors of x , if y were an ancestor of p (meaning y depends on p), it being in the cache would violate p 's status as an eviction candidate. Thus, such an unmarked y cannot exist in the cache.

Therefore, every page of C' currently in the cache is marked. By definition, C' is a marked chain. $\square$

Let u denote the total number of unmarked chains at the beginning of the phase. By definition, an unmarked chain must possess at least one unmarked page currently residing in the cache. Because the cache capacity is exactly k , there can be at most k such chains. Since there are l chains in total, we have $u \leq \min(k, l)$.

Furthermore, a marked chain never becomes unmarked during a phase. Let $C_1, C_2, \dots, C_u$ denote these initially unmarked chains, indexed in the exact order they become marked during the phase. Let t_j denote the number of coins residing on chain C_j at the exact moment it becomes marked.

Claim 5 (Expected Coins on Marked Chains). *When the j -th initially unmarked chain (C_j) becomes marked, it possesses in expectation at most $\frac{m}{u-j+1}$ coins. That is:*

$$E[t_j] \leq \frac{m}{u-j+1}$$

Proof. Consider the exact moment just before C_j becomes fully marked. At this moment, exactly $j-1$ initially unmarked chains have already been marked, meaning there are exactly $u-j+1$ unmarked chains remaining in the system (including C_j).

By Claim 4, we only remove coins from marked chains. This means there may be coins residing on chains that are already marked, but this is only in our favor, as it leaves fewer coins distributed among the remaining unmarked chains.

By Claim 3, the total number of coins in the entire system at any given moment is bounded by m . Since the algorithm distributes new coins uniformly at random among all currently unmarked chains, by symmetry, the expected number of coins residing on any specific unmarked chain is bounded by the maximum total coins in the system divided by the number of unmarked chains, $u-j+1$. Therefore:

$$E[t_j] \leq \frac{m}{u-j+1}$$

$\square$

Theorem 2. *The randomized dependency-aware marking algorithm is $O(\min(\log k, \log l))$ -competitive.*

Proof. We are interested in bounding the number of cache misses on old pages. An old page incurs a miss if and only if it was evicted and subsequently refetched into the cache.

By our accounting rules, an old page takes charge of a coin exactly when it is evicted, and this coin is discarded exactly when the page is refetched. Therefore, the total number of misses on old pages provides an upper bound that is exactly equal to the total number of coins discarded during the phase.

By Claim 4, coins are exclusively discarded from a chain once it becomes marked. Because marked chains do not receive new coins (as evictions only target unmarked chains), the maximum number of coins that can be discarded from chain C_j is bounded by t_j , the number of coins it possessed at the exact moment it became marked. Summing these expectations over all u initially unmarked chains provides an upper bound on the total expected cost for old pages:

$$\begin{aligned} E[\text{misses on old pages}] &\leq \sum_{j=1}^u E[t_j] \\ &\leq \sum_{j=1}^u \frac{m}{u-j+1} \quad (\text{by Claim 5}) \\ &= m \sum_{i=1}^u \frac{1}{i} \\ &= O(m \log u) \\ &= O(m \min(\log k, \log l)) \quad (\text{since } u \leq \min(k, l)) \end{aligned}$$

This bounds the expected cost of fetching old pages. Furthermore, the algorithm incurs exactly m misses to fetch the m new pages requested during the phase. Thus, the total expected cost of the algorithm for the phase is $m + O(m \min(\log k, \log l)) = O(m \min(\log k, \log l))$.

By Lemma 1, OPT incurs an amortized cost of at least $m/2$ misses per phase. Therefore, our randomized algorithm yields an overall expected competitive ratio of $O(\min(\log k, \log l))$. $\square$

References

- [1] Julien Dallot et al. “Dependency-aware online caching”. In: *IEEE INFOCOM 2024-IEEE Conference on Computer Communications*. IEEE. 2024, pp. 871–880.