

Dependency Paging

Tal Zohar

May 21, 2026

1 The Dependency Paging Problem

There are n pages and a cache of size k . Additionally, we are provided with a dependency DAG G . Let $G[p]$ be the set of dependencies of page p . A valid solution must ensure that if page p is in cache, then all pages in $G[p]$ are also in the cache. Let l denote the width of G - the maximum antichain length.

We analyze the problem assuming pages are requested one by one, such that every request arrives only after all of its children (dependencies) have been requested. This topological request order simplifies the algorithm description and proofs without loss of generality.

2 The Unweighted Case - Marking Algorithms and Phase Structure

To analyze both the deterministic and randomized algorithms for the unweighted problem, we utilize the classical framework of *marking algorithms*, adapted to respect dependency constraints.

A marking algorithm processes a request sequence in a series of disjoint phases. At the beginning of a phase, all pages are considered “unmarked.” As requests arrive, the requested pages (and their dependencies) are “marked.” The fundamental rule of a marking algorithm is that it **never evicts a marked page**. When the cache is full and an eviction is required, the algorithm is forced to choose an unmarked page. A phase naturally terminates when an eviction is needed but the cache is entirely full of marked pages. At this point, all marks are cleared, and a new phase begins.

Consequently, the sequence is divided into phases containing exactly k distinct pages, where k represents the size of the cache. Let m denote the number of new pages requested in a phase (pages that were not in the cache at the start of the phase).

Lemma 1 (OPT Lower Bound). *The optimal offline algorithm (OPT) incurs an amortized cost of at least $m/2$ misses per phase.*

Proof. Because OPT has a cache capacity of exactly k , and each phase requests exactly k distinct pages including m entirely new pages, OPT must incur at least m misses across the current and previous phases combined. Thus, its amortized cost per phase is at least $m/2$. $\square$

By Lemma 1, any algorithm that incurs $O(m \cdot f(l))$ misses per phase is $O(f(l))$ -competitive.

We partition the DAG G into l disjoint strongly-ordered chains $C_1, \dots, C_l$, which is possible by Dilworth’s theorem.

Definition: A page is considered **marked** if it has been requested during the current phase.

Definition: A chain is considered **marked** if there is no unmarked page of the chain currently residing in the cache (i.e., all chain pages in the cache are marked).

Definition: A page in the cache is an **eviction candidate** if it is a maximal unmarked cached page (meaning it is unmarked, and no other page currently in the cache is dependent on it). Because requests follow a topological order, these are the only pages that can be safely evicted without violating dependency constraints.

Observation 1. Equivalence of Phase Boundaries: *By definition, a phase ends when all k pages in the cache are marked. Because the cache capacity is exactly k , this is the exact moment when no unmarked pages remain in the cache, and consequently, all l chains become marked. Therefore, a phase runs exactly as long as there are valid eviction candidates in the cache.*

Observation 2. *Within each chain, the marked pages reside at the bottom, while the evicted pages reside at the top.*

Observation 3. *Once a chain becomes marked, it never becomes unmarked during the phase, since pages fetched mid-phase are immediately marked upon entry.*

3 A Deterministic $O(\min(k, l))$ -Competitive Algorithm for the Unweighted Case

Standard LRU preserves dependency constraints and is $O(k)$ -competitive, as shown in [2]. By a tighter analysis, we can show that a generic dependency-aware marking algorithm is $O(\min(k, l))$ -competitive.

Claim 1. *Standard LRU is a valid implementation of the generic deterministic dependency-aware marking algorithm.*

Algorithm 1 describes this generic deterministic strategy. Whenever the cache is full, it simply evicts any valid eviction candidate.

Algorithm 1 Deterministic Marking Algorithm for Dependency Paging

```

1: Initialize: Cache  $S \leftarrow \emptyset$ , Marked set  $M \leftarrow \emptyset$ 
2: for each requested page  $p$  in the sequence do
3:   if  $|M \cup \{p\}| > k$  then                                 $\triangleright$  Phase ends when we attempt to mark beyond capacity
4:     End current phase, start new phase
5:      $M \leftarrow \emptyset$                                         $\triangleright$  Unmark all pages
6:     Mark all pages in  $S$  that are needed for  $p$ 
7:   end if
8:    $M \leftarrow M \cup \{p\} \cup G[p]$                              $\triangleright$  Mark the requested page and its dependencies
9:   if  $|S \cup \{p\}| > k$  then                                     $\triangleright$  Cache is full, eviction required
10:    Find an eviction candidate  $q \in S$ 
11:     $S \leftarrow S \setminus \{q\}$                                  $\triangleright$  Evict page  $q$ 
12:  end if
13:  if  $p \notin S$  then
14:     $S \leftarrow S \cup \{p\}$                                      $\triangleright$  Fetch the requested page into the cache
15:  end if
16: end for

```

To prove the competitiveness, we bound the number of misses on “old pages” (pages that were in the cache at the start of the phase).

Lemma 2 (Bound on Displaced Pages). *At any given moment during a phase, the total number of old pages that are currently evicted from the cache is bounded by m .*

Proof. Because the cache maintains a fixed capacity of k , every currently evicted old page must have been perfectly replaced by a “new page” requested during this phase. Since at most m new pages are requested in the entire phase, the cache can physically hold at most m new pages at any given time. Thus, the total number of evicted old pages is bounded by m . $\square$

Lemma 3 (Chain Marking Upon Miss). *Let p be an old page belonging to chain C_j that was evicted during the current phase. At the exact moment p is subsequently requested and brought back into the cache, C_j becomes a marked chain.*

Proof. Because C_j is a totally ordered chain, any other page $y \in C_j$ must either be a dependency of p , or y must depend on p .

- If y is a dependency of p , the algorithm guarantees that when p is marked, all of its dependencies currently in the cache are also marked. Thus, y is marked.
- If y depends on p , then y requires p to be in the cache. When p was originally evicted, it was an eviction candidate, meaning no cached page depended on it. Since p has remained evicted until now, no page depending on p (including y) could have been requested or fetched into the cache, as doing so would require fetching p first. Thus, y is currently not in the cache.

Therefore, every page of C_j currently in the cache is marked, making C_j a marked chain. $\square$

Theorem 1. *The generic deterministic marking algorithm is $O(\min(k, l))$ -competitive.*

Proof. We are interested in bounding the total number of misses on old pages. An old page incurs a miss if and only if it is evicted and subsequently refetched.

First, observe that once an old page is refetched, it becomes marked. By the fundamental rule of marking algorithms, marked pages are never evicted. Thus, each of the at most k old pages originally in the cache can incur a maximum of one miss per phase. This bounds the total number of misses on old pages strictly by k .

Second, we can bound the misses using the chain structure. By Lemma 3, the very first time an evicted old page from C_j is refetched, C_j becomes a marked chain. By Observation 3, a marked chain never becomes unmarked again for the remainder of the phase.

Because the algorithm only evicts unmarked pages (eviction candidates), no page from C_j can be evicted once C_j becomes marked. Therefore, all evictions from C_j must occur *before* C_j suffers its first miss.

By Lemma 2, the maximum number of pages from C_j that can be simultaneously evicted at any given moment is bounded by the total number of evicted old pages in the entire system, which is m . Since C_j cannot be evicted from after its first miss, the maximum number of misses C_j can incur is exactly the number of its pages that were evicted prior to this moment, which is at most m .

Summing this bound over all l chains, the total number of misses on old pages is also bounded by $m \cdot l$.

Combining these two independent bounds, the total number of misses on old pages is at most $\min(k, m \cdot l) \leq m \cdot \min(k, l)$ (since $m \geq 1$).

The algorithm also incurs exactly m misses to fetch the m new pages. Thus, the total cost for the phase is at most $m + m \cdot \min(k, l) = m(1 + \min(k, l))$. By Lemma 1, OPT incurs an amortized cost of at least $m/2$ per phase. Therefore, the generic deterministic marking algorithm achieves an overall competitive ratio of $O(\min(k, l))$. $\square$

4 A Randomized $O(\min(\log k, \log l))$ -Competitive Algorithm for the Unweighted Case

The paper [2] provides an expected $O(\min(\log k, \log l))$ -competitive randomized algorithm. Here, we provide an alternative, simpler analysis to establish this exact $O(\min(\log k, \log l))$ competitive ratio directly.

We present a randomized algorithm (Algorithm 2) that restricts how the eviction candidates are chosen. By ensuring evictions are distributed symmetrically across the chains, we prevent the adversary from stacking m misses on a single chain.

4.1 Blocked Chains and Algorithm Description

We reuse the chains $C_1, \dots, C_l$ and the definition of an eviction candidate. To formalize the randomized selection, we define how a chain specifies its desired eviction target.

Definition: For any unmarked chain, its **next eviction candidate** is defined as its highest unmarked page currently residing in the cache. Note that a chain's next eviction candidate might not be a global *eviction candidate*, as there could be cached pages *outside* the chain that depend on it.

We introduce the concept of a “blocked chain” to systematically determine how actual eviction candidates are chosen.

Definition: An unmarked chain is considered **blocked** on a page p if the following two conditions are met:

1. p is an eviction candidate.
2. p is an ancestor of (or equal to) the chain's next eviction candidate.

Claim 2. *Every unmarked chain must be blocked by some page (can be multiple).*

Proof. Let $\tilde{C}$ be an unmarked chain. By definition, it has a next eviction candidate $x \in \tilde{C}$ currently residing in the cache. If we trace the dependency graph upwards from x to its ancestors currently in the

cache, we will eventually reach a maximal unmarked cached page q (one with no cached dependents). This page q is an eviction candidate. Because q is either equal to x or an ancestor of x , the chain $\tilde{C}$ is blocked on q . $\square$

Eviction Rule: When we need to evict a page, we draw an unmarked chain $\tilde{C}$ uniformly at random. We select some page q such that $\tilde{C}$ is blocked on q (by Claim 2, we are guaranteed such a page exists), and we evict q .

Algorithm 2 Randomized Marking Algorithm for Dependency Paging

```

1: Initialize: Cache  $S \leftarrow \emptyset$ , Marked set  $M \leftarrow \emptyset$ 
2: Partition  $G$  into  $l$  disjoint chains  $C_1, \dots, C_l$  ▷ via Dilworth's Theorem
3: for each requested page  $p$  in the sequence do
4:   if  $|M \cup \{p\}| > k$  then ▷ Phase ends when we attempt to mark beyond capacity
5:     End current phase, start new phase
6:      $M \leftarrow \emptyset$  ▷ Unmark all pages
7:     Mark all pages in  $S$  that are needed for  $p$ 
8:   end if
9:    $M \leftarrow M \cup \{p\} \cup G[p]$  ▷ Mark the requested page and its dependencies
10:  if  $|S \cup \{p\}| > k$  then ▷ Cache is full, eviction required
11:    Draw an unmarked chain  $\tilde{C}$  uniformly at random
12:    Select a page  $q \in S$  such that  $\tilde{C}$  is blocked on  $q$ 
13:    ▷ ( $q$  is an eviction candidate, and ancestor of  $\tilde{C}$ 's next eviction candidate)
14:     $S \leftarrow S \setminus \{q\}$  ▷ Evict page  $q$ 
15:  end if
16:  if  $p \notin S$  then
17:     $S \leftarrow S \cup \{p\}$  ▷ Fetch the requested page into the cache
18:  end if
19: end for

```

4.2 Conceptual Accounting Scheme and Analysis

To bound the expected number of cache misses, we introduce a conceptual accounting scheme using “coins”. It is important to note that these coins are strictly an analytical tool used for the proof; the algorithm does not track or utilize them during its execution.

Coin Distribution Rules:

- **Giving Coins:** Whenever the algorithm evicts a page q by randomly selecting an unmarked chain $\tilde{C}$, we conceptually give one coin to chain $\tilde{C}$. We define the evicted page q to be the page “in charge” of this coin.
- **Removing Coins:** If a page q that is in charge of a coin is subsequently requested (and thus marked), we remove its associated coin from its chain. This scenario occurs exclusively when an evicted old page is refetched into the cache.

Based on this accounting scheme and the structural properties established in Section 2, we establish the following claims:

Claim 3 (Bound on Total Coins). *At any given moment, the total number of coins in the system is bounded by m .*

Proof. By our accounting rules, a coin corresponds exactly to an old page that is currently evicted. By Lemma 2, the number of currently evicted old pages is bounded by m . Thus, the total number of coins is bounded by m . $\square$

Claim 4 (Coin Removal). *We only remove coins from marked chains.*

Proof. A coin is removed from a chain C' exactly when the page p in charge of it is requested. We must show that at this moment, C' is a marked chain (i.e., it has no unmarked pages in the cache).

When p was originally evicted and took charge of the coin, C' was blocked on p . By definition, this means C' had a next eviction candidate x (its highest unmarked cached page at that time), and p was an ancestor of x (or equal to x).

Fast forward to the moment p is requested and the coin is removed. Because p is requested, p is immediately marked. The algorithm guarantees that when a page is marked, all its dependencies currently in the cache are also marked. Because p is an ancestor of x , x is a dependency of p . Thus, x becomes marked.

Could there be some other page $y \in C'$ in the cache that is unmarked? Because C' is a totally ordered chain of dependencies, y must be either a dependency of x , or x must be a dependency of y .

- If y is a dependency of x , then by transitivity, y is a dependency of p . For the same reason x is marked, y must also be marked.
- If x is a dependency of y , then y is an ancestor of x . However, recall that when p was evicted, it was an eviction candidate, meaning it had *no* cached dependents (no cached pages depending on it). Because C' is a totally ordered chain, y and p must be ordered. Since both y and p are ancestors of x , if y were an ancestor of p (meaning y depends on p), it being in the cache would violate p 's status as an eviction candidate. Thus, such an unmarked y cannot exist in the cache.

Therefore, every page of C' currently in the cache is marked. By definition, C' is a marked chain. $\square$

Let u denote the total number of unmarked chains at the beginning of the phase. By definition, an unmarked chain must possess at least one unmarked page currently residing in the cache. Because the cache capacity is exactly k , there can be at most k such chains. Since there are l chains in total, we have $u \leq \min(k, l)$.

Furthermore, by Observation 3, a marked chain never becomes unmarked during a phase. Let $C_1, C_2, \dots, C_u$ denote these initially unmarked chains, indexed in the exact order they become marked during the phase. Let t_j denote the number of coins residing on chain C_j at the exact moment it becomes marked.

Claim 5 (Expected Coins on Marked Chains). *When the j -th initially unmarked chain (C_j) becomes marked, it possesses in expectation at most $\frac{m}{u-j+1}$ coins. That is:*

$$E[t_j] \leq \frac{m}{u-j+1}$$

Proof. Consider the exact moment just before C_j becomes fully marked. At this moment, exactly $j-1$ initially unmarked chains have already been marked, meaning there are exactly $u-j+1$ unmarked chains remaining in the system (including C_j).

By Claim 4, we only remove coins from marked chains. This means there may be coins residing on chains that are already marked, but this is only in our favor, as it leaves fewer coins distributed among the remaining unmarked chains.

By Claim 3, the total number of coins in the entire system at any given moment is bounded by m . Since the algorithm distributes new coins uniformly at random among all currently unmarked chains, by symmetry, the expected number of coins residing on any specific unmarked chain is bounded by the maximum total coins in the system divided by the number of unmarked chains, $u-j+1$. Therefore:

$$E[t_j] \leq \frac{m}{u-j+1}$$

$\square$

Theorem 2. *The randomized dependency-aware marking algorithm is $O(\min(\log k, \log l))$ -competitive.*

Proof. We are interested in bounding the number of cache misses on old pages. An old page incurs a miss if and only if it was evicted and subsequently refetched into the cache.

By our accounting rules, an old page takes charge of a coin exactly when it is evicted, and this coin is discarded exactly when the page is refetched. Therefore, the total number of misses on old pages provides an upper bound that is exactly equal to the total number of coins discarded during the phase.

By Claim 4, coins are exclusively discarded from a chain once it becomes marked. Because marked chains do not receive new coins (as evictions only target unmarked chains), the maximum number of coins that can be discarded from chain C_j is bounded by t_j , the number of coins it possessed at the exact

moment it became marked. Summing these expectations over all u initially unmarked chains provides an upper bound on the total expected cost for old pages:

$$\begin{aligned}
E[\text{misses on old pages}] &\leq \sum_{j=1}^u E[t_j] \\
&\leq \sum_{j=1}^u \frac{m}{u-j+1} \quad (\text{by Claim 5}) \\
&= m \sum_{i=1}^u \frac{1}{i} \\
&= O(m \log u) \\
&= O(m \min(\log k, \log l)) \quad (\text{since } u \leq \min(k, l))
\end{aligned}$$

This bounds the expected cost of fetching old pages. Furthermore, the algorithm incurs exactly m misses to fetch the m new pages requested during the phase. Thus, the total expected cost of the algorithm for the phase is $m + O(m \min(\log k, \log l)) = O(m \min(\log k, \log l))$.

By Lemma 1, OPT incurs an amortized cost of at least $m/2$ misses per phase. Therefore, our randomized algorithm yields an overall expected competitive ratio of $O(\min(\log k, \log l))$. $\square$

5 Dependency Paging with Rejections

In the standard dependency paging model, every requested page must be fetched into the cache. We now consider a variant where the algorithm is given the choice to either fetch a requested page or *bypass* it, incurring a penalty. In standard paging, it is a well-established result that algorithms designed for the non-bypass setting can be systematically converted into algorithms that support bypassing, a framework formalized under the problem of *paging with rejections* [4].

However, introducing bypass decisions to the dependency paging problem adds significant complexity that does not exist in standard paging. Because pages are constrained by a dependency DAG, bypassing a page inherently means that its required dependencies also do not need to be fetched to satisfy that specific request. Due to this cascading effect on structural requirements, the known conversion techniques from standard paging cannot be directly applied to our model.

Recent work has attempted to bridge this gap. The authors in [2] adapted their algorithm to support dependency paging with bypasses (where the bypass cost is uniformly fixed at 1), achieving an $O(\sqrt{k \min(\log k, \log l)})$ -competitive ratio. Another related study on online tree caching [1] provides an algorithm for a generalized problem where bypassing incurs a fixed cost α . They establish a competitive ratio of $O(k \cdot h(T))$, where $h(T)$ is the height of the dependency tree. However, their analysis relies heavily on the underlying topology and only applies when the dependency structure forms a tree.

In this section, our objective is not to improve upon the competitive ratio established by [2], but rather to formally extend their result to an even more generalized problem: arbitrary *paging with rejections* [4]. In this generalized setting, every individual request has its own distinct bypass cost $b \in \mathbb{R}^+$. Note that this cost is associated with the request itself and does not need to be inherently tied to the requested page. Remarkably, we demonstrate that our adaptation accommodates this flexibility and maintains the state-of-the-art $O(\sqrt{k \min(\log k, \log l)})$ competitive ratio while navigating the structural challenges of arbitrary bypasses in general DAGs.

The Cost of Bypasses and the Failure of Top-Down Assumptions

Formally, the input sequence consists of requests denoted by the tuple (p, b) , where p is the requested page and $b \in \mathbb{R}^+$ is the penalty incurred if the algorithm chooses to bypass this specific request. In standard paging with rejections, one can usually assume without loss of generality that the bypass cost is no greater than the cost of a cache fetch (i.e., $b \leq 1$ in the unweighted case). This is because bypassing a request for a cost greater than fetching it is trivially suboptimal [4].

However, in dependency paging, we cannot make this assumption. Fetching a page p strictly requires all its dependencies $G[p]$ to be in the cache. If multiple dependencies of p are currently missing, the true cost of fetching p is the cost of fetching p *plus* the cost of fetching the entire missing dependency chain.

Therefore, an adversary might present a request with a bypass cost of $b > 1$, and it might still be cheaper for the algorithm to bypass the request than to pay for a massive chain of prerequisite fetches.

To understand how to adapt our algorithm, we must first examine where our previous analysis fails in the presence of bypasses. In the non-bypass setting, we relied heavily on a topological assumption: requests arrive such that when a page p is requested, all of its dependencies are already marked.

In the bypass setting, this assumption collapses because algorithms—specifically the optimal offline algorithm (OPT)—now have the freedom to skip requests. In the standard model, every valid algorithm is forced to serve a request for a dependency $q \in G[p]$, guaranteeing it is marked before p is ever requested. However, when bypass decisions are introduced, OPT (or our online algorithm) can strategically choose to bypass a prior request for q . Consequently, q is neither fetched nor marked. If p is subsequently requested and the algorithm decides to fetch it, it is forced to simultaneously fetch p and all of its unmarked dependencies. This bulk-fetching is catastrophic for the competitive ratio. The core of our non-bypass analysis rested on the invariant that a single fetch causes at most one eviction, bounding the total number of “holes” by m . If a single request triggers multiple cascading fetches, it causes a burst of simultaneous evictions, inflating the holes well beyond m and completely destroying the structural bounds.

The Fractional Budget-Pouring Algorithm

To restore the topological invariant, we must ensure that pages are only fetched one at a time, strictly after their dependencies have been satisfied. We achieve this by treating the bypass cost scaled by a parameter $\beta > 0$ as a continuous liquid budget. Let $\prec$ be an arbitrary but strictly fixed total order over the pages of G .

Let $w_q \in [0, 1]$ be a fractional counter for each page $q \in G$, initialized to 0. Upon receiving a request (p, b) , let U_p denote the set of all currently unmarked pages in $G[p] \cup \{p\}$. If U_p is empty, all required pages are already marked, so we fetch and serve p for free.

If U_p is not empty, we define $S(p) \subseteq U_p$ as the set of unmarked pages whose dependencies are *entirely marked*. We consistently identify the single page $q \in S(p)$ that possesses the highest index according to our fixed order $\prec$. We continuously pour the scaled budget $\beta \cdot b$ exclusively into w_q . Crucially, to preserve synchronization with the optimal offline cost, we pour the budget into this page regardless of whether it is already residing in the cache; it is only marked and considered fully satisfied when its counter w_q reaches 1.

As soon as w_q reaches 1, the page conceptually prepares to be marked. We first check for a phase termination: if marking q would exceed the cache capacity k , the current phase ends. All marks are cleared ($M \leftarrow \emptyset$), and all fractional counters are reset to 0. We then recompute U_p and simply *continue* the budget-pouring loop from scratch.

If no phase reset is triggered (or after a reset occurs and q successfully reaches $w_q = 1$ again), q is marked and removed from U_p . If fetching q into the cache requires an eviction, we execute exactly one randomized chain eviction. This sequential pouring process continues dynamically until either the scaled budget $\beta \cdot b$ is completely exhausted (in which case the request for p is bypassed), or U_p becomes empty, meaning all required pages have been conceptually fetched and marked.

The Online Reduction and Analysis

This budget-pouring algorithm can be formally viewed as an online reduction that takes the original input sequence I (with bypasses) and converts it into a virtual sequence I' of standard (non-bypass) requests.

Whenever a page q 's counter reaches 1 during the pouring process, we append a virtual request for q to I' . Let the t -th request in I be denoted as req_t . Processing req_t generates a (possibly empty) sequence of virtual requests $req'_{t,1}, req'_{t,2}, \dots$ in I' . Crucially, because we restrict the target selection strictly to the set $S(p)$ —pages whose dependencies are *already marked*—the virtual sequence I' is constructed in such a way that the topological assumption is perfectly upheld. In I' , a page is only requested when its dependencies are guaranteed to be marked.

To analyze the competitive ratio on this reduced sequence, we redefine m_i precisely as the number of *new marked pages* in phase i relative to the previous phase. Formally, $m_i = |M_i \setminus M_{i-1}|$, where M_i is the set of marked pages at the end of phase i , and M_{i-1} is the set of marked pages at the end of phase $i - 1$. By this definition, m_i exactly captures the number of new marks introduced by the virtual sequence I' during phase i .

Algorithm 3 Randomized Budget-Pouring for Paging with Rejections

```

1: Initialize: Cache  $S \leftarrow \emptyset$ , Marked set  $M \leftarrow \emptyset$ 
2: Partition  $G$  into  $l$  disjoint chains  $C_1, \dots, C_l$  ▷ via Dilworth's Theorem
3: Fix an arbitrary strict total order  $\prec$  on the pages of  $G$ 
4: Initialize  $w_y \leftarrow 0$  for all  $y \in G$ 
5: Let  $\beta > 0$  be a designated scaling parameter
6: for each request  $(p, b)$  in the sequence do
7:    $U_p \leftarrow \{y \in G[p] \cup \{p\} \mid y \notin M\}$  ▷ Unmarked required pages
8:   if  $U_p = \emptyset$  then
9:     Fetch  $p$  if  $p \notin S$ 
10:    continue ▷ Cost is 0, dependencies already marked
11:  end if
12:  Let  $B_{rem} \leftarrow \beta \cdot b$  be the continuous budget
13:  while  $B_{rem} > 0$  and  $U_p \neq \emptyset$  do
14:     $S(p) \leftarrow \{y \in U_p \mid \text{all dependencies of } y \text{ are in } M\}$ 
15:    Let  $q$  be the page in  $S(p)$  with the highest index according to  $\prec$ 
16:    Pour budget continuously into  $w_q$ 
17:    Decrease  $B_{rem}$  by the total amount poured
18:    if  $w_q$  reaches 1 then
19:      if  $|M \cup \{q\}| > k$  then ▷ Phase ends when marking beyond capacity
20:        End current phase, start new phase
21:         $M \leftarrow \emptyset$  ▷ Unmark all pages
22:         $w_y \leftarrow 0$  for all  $y \in G$  ▷ Reset counters for the new phase
23:         $U_p \leftarrow \{y \in G[p] \cup \{p\} \mid y \notin M\}$  ▷ Recompute missing pages
24:        continue ▷ Continue pouring from scratch
25:      end if
26:       $M \leftarrow M \cup \{q\}$  ▷ Mark the page
27:      if  $|S \cup \{q\}| > k$  then ▷ Cache is full, exactly one eviction required
28:        Draw an unmarked chain  $\tilde{C}$  uniformly at random
29:        Select a page  $x \in S$  such that  $\tilde{C}$  is blocked on  $x$ 
30:         $S \leftarrow S \setminus \{x\}$  ▷ Evict page  $x$ 
31:      end if
32:      if  $q \notin S$  then
33:         $S \leftarrow S \cup \{q\}$  ▷ Fetch the page into the cache
34:      end if
35:       $U_p \leftarrow U_p \setminus \{q\}$  ▷ Page is satisfied, remove from missing set
36:    end if
37:  end while
38:  if  $U_p = \emptyset$  then
39:    Request  $p$  is fully fetched and marked
40:  else
41:    Request  $p$  is bypassed
42:  end if
43: end for

```

To bound the total cost incurred by our algorithm (ALG), we decompose its cost into two components: the cost of fetches (cache misses) and the cost of bypasses (the total penalty incurred).

Lemma 4 (Single Active Fractional Page). *At any given time during the execution of the algorithm, there exists at most one page $q \in G$ such that $0 < w_q < 1$.*

Proof. The subset of eligible pages $S(p)$ only changes when a page's counter reaches 1 (resulting in it being marked and removed from U_p) or when a new request arrives. Between these discrete events, the set $S(p)$ is strictly static, meaning the page $q \in S(p)$ with the highest index according to the global order $\prec$ remains uniquely constant. Thus, the algorithm pours budget exclusively into a single counter until it caps at 1. Since all counters initialize at 0 and are only incremented one strictly defined target at a time, exactly one page can possess a strictly fractional budget at any moment. $\square$

Lemma 5 (Phase Bound on Fetch Cost). *In any phase i , the expected fetch cost incurred by ALG is bounded by $O(m_i \min(\log k, \log l))$.*

Proof. The budget-pouring mechanism constructs the virtual sequence I' such that a virtual request for a page q is issued only when all of its dependencies $G[q]$ are already marked. Consequently, I' strictly adheres to the topological ordering constraint. Running Algorithm 3 on I' generates exactly m_i new marked pages in phase i . Because the topological structure and the randomized eviction strategy remain identical, the core analysis from Section 4 applies directly to each phase. Thus, the expected number of cache misses (fetches) generated by ALG in phase i is exactly $O(m_i \min(\log k, \log l))$. $\square$

To cleanly analyze the bypass penalty when a single request (p, b) might trigger a cache capacity limit and straddle a phase boundary mid-pour, we measure the bypass cost continuously. We define the **true bypass cost** of phase i , denoted B_i , as the exact volume of unscaled penalties b consumed strictly between the start and the end of phase i . Because the algorithm scales b by β , the total volume of fractional budget actually poured during phase i is exactly $\beta \cdot B_i$. Summing B_i across all phases strictly upper-bounds ALG 's actual total bypass cost.

Lemma 6 (Phase Bound on Bypass Cost). *In any phase i , the true bypass cost B_i incurred by ALG is strictly bounded by $\frac{1}{\beta}$ times the total number of marks in that phase plus $\frac{1}{\beta}$. That is, $B_i < \frac{1}{\beta}(|M_i| + 1)$.*

Proof. By Lemma 4, budget is poured into exactly one active page at a time. Within the temporal window of phase i , every full 1 unit of poured budget converts exactly into one page marking event, regardless of whether the target page was already in the cache. The phase concludes (either by hitting the capacity limit k and triggering a reset, or by reaching the end of the sequence) with at most one page holding a residual fractional budget $0 \leq w_q < 1$. Because this residual fraction is strictly less than 1, the total volume of budget poured strictly during phase i (which equals $\beta \cdot B_i$) is exactly bounded by the number of marks made ($|M_i|$) plus this residual fraction. Thus, $\beta \cdot B_i < |M_i| + 1$, which implies $B_i < \frac{1}{\beta}(|M_i| + 1)$. $\square$

Lemma 7 (Total Upper Bound for ALG). *The total expected cost incurred by the budget-pouring algorithm over the entire sequence is bounded by $\sum_i \left(O(m_i \min(\log k, \log l)) + \frac{k}{\beta} \right)$.*

Proof. The total cost paid by ALG is strictly bounded by the sum of its expected fetch costs and its true bypass costs across all phases. By Lemma 5, the expected fetch cost in phase i is $O(m_i \min(\log k, \log l))$. By Lemma 6, the true bypass cost in phase i satisfies $B_i < \frac{1}{\beta}(|M_i| + 1)$. Since exactly k pages are marked in every fully completed phase (except possibly the final one), $|M_i| \leq k$. Combining these bounds yields the total cost per phase:

$$E[\text{Total Cost}_i] \leq E[\text{Fetch Cost}_i] + B_i \leq O(m_i \min(\log k, \log l)) + \frac{k+1}{\beta} = O(m_i \min(\log k, \log l)) + \frac{k}{\beta}$$

Summing this independent bound across all phases establishes the exact mathematical upper bound on the online algorithm's total expected cost:

$$E[\text{Total Cost}] \leq \sum_i \left(O(m_i \min(\log k, \log l)) + \frac{k}{\beta} \right)$$

$\square$

To evaluate the competitive ratio, we bound the optimal offline cost against this redefined parameter:

Lemma 8 (Global Bypass OPT Lower Bound). *Let $m_i = |M_i \setminus M_{i-1}|$ be the number of new marked pages in phase i . Under the algorithm where budget is poured into targets regardless of their cache status, the optimal offline algorithm (OPT) incurs a total cost of at least $\Omega\left(\min(1, \frac{1}{\beta}) \sum_i m_i\right)$ over the entire sequence.*

Proof. To establish the lower bound, we analyze the cost incurred by OPT over pairs of consecutive phases. Consider any continuous temporal window spanning two consecutive phases, $i-1$ and i .

Let $M = M_{i-1} \cup M_i$ be the set of pages marked by our online algorithm (ALG) during this window. By definition, $|M_{i-1}| = k$ and $m_i = |M_i \setminus M_{i-1}|$, meaning the total number of distinct pages in M is exactly $k + m_i$. Let $X \subseteq M$ denote the subset of marked pages that OPT *never holds* in its cache at any point during these two phases.

First, we bound OPT's fetch cost. Over the duration of phases $i-1$ and i , the maximum number of distinct pages OPT can hold in its cache is bounded by its initial cache capacity at the start of phase $i-1$ (which is k) plus the total number of new pages it fetches. Let $f_{OPT}^{(i-1,i)}$ denote the number of fetches OPT makes during this window. Because OPT never caches X , it can only hold a maximum of $k + f_{OPT}^{(i-1,i)}$ pages from the set $M \setminus X$. By the fundamental requirement of the sequences, OPT must eventually hold all required pages not bypassed. Since it bypassed X , it must fetch the remainder. Therefore:

$$k + f_{OPT}^{(i-1,i)} \geq |M \setminus X| = k + m_i - |X| \implies f_{OPT}^{(i-1,i)} \geq m_i - |X|$$

Since the number of fetches cannot be negative, OPT's fetch cost during this time window is at least $\max(0, m_i - |X|)$.

Second, we bound OPT's bypass cost using the fractional budget mechanics. Let $\mathcal{B}_{OPT}$ be the set of all requests bypassed by OPT strictly during the time window of phases $i-1$ and i . Let $w_{t,q}$ denote the exact amount of continuous budget poured by ALG into the counter of page q while processing request t .

For every request $t \in \mathcal{B}_{OPT}$ with penalty b_t , ALG pours a total continuous budget that never exceeds $\beta \cdot b_t$. Therefore, β times OPT's true bypass cost is bounded by the summation of the budget ALG distributed across these specific requests:

$$\beta \cdot \text{Bypass Cost}_{OPT}^{(i-1,i)} = \sum_{t \in \mathcal{B}_{OPT}} \beta \cdot b_t \geq \sum_{t \in \mathcal{B}_{OPT}} \left(\sum_{q \in X} w_{t,q} \right)$$

By swapping the order of summation, we strictly isolate the fractional budget deposited into the un-cached pages X :

$$\sum_{t \in \mathcal{B}_{OPT}} \sum_{q \in X} w_{t,q} = \sum_{q \in X} \left(\sum_{t \in \mathcal{B}_{OPT}} w_{t,q} \right)$$

Consider any page $q \in X$. Because $q \in M_{i-1} \cup M_i$, ALG marked q precisely during phase $i-1$ or phase i . Crucially, because ALG strictly pours budget into target pages regardless of whether they are already in the cache, ALG was forced to pour exactly 1 total unit of budget into q to successfully mark it. Because all fractional counters are reset to 0 at phase boundaries, this 1 unit of budget was poured exclusively within the boundaries of the phase in which q was marked. Thus, the 1 unit of budget was strictly poured within our two-phase temporal window.

Every time ALG pours budget into q while processing some request t , q is an inherently required dependency for t . Because $q \in X$, OPT does *not* have q in its cache at any point during these phases. Consequently, OPT is physically incapable of serving request t and *must* bypass it (meaning $t \in \mathcal{B}_{OPT}$).

Therefore, *every single request* that contributes to $w_{t,q}$ for $q \in X$ is inherently a request bypassed by OPT. The inner sum $\sum_{t \in \mathcal{B}_{OPT}} w_{t,q}$ thus perfectly captures the entirety of the 1 unit of budget required to mark q . This reduces the inequality to:

$$\beta \cdot \text{Bypass Cost}_{OPT}^{(i-1,i)} \geq \sum_{q \in X} 1 = |X| \implies \text{Bypass Cost}_{OPT}^{(i-1,i)} \geq \frac{|X|}{\beta}$$

Combining the bounds, the total cost incurred by OPT during phases $i-1$ and i is:

$$\text{Total Cost}_{OPT}^{(i-1,i)} \geq \text{Fetch Cost}_{OPT}^{(i-1,i)} + \text{Bypass Cost}_{OPT}^{(i-1,i)} \geq \max(0, m_i - |X|) + \frac{|X|}{\beta}$$

Since $0 \leq |X| \leq m_i$, the minimum value of this combination depends on the magnitude of β . It is minimized either at $|X| = 0$ or $|X| = m_i$, yielding a cost bounded by $\min(1, \frac{1}{\beta})m_i$.

Since OPT pays at least $\min(1, \frac{1}{\beta})m_i$ for every overlapping pair of consecutive phases $(i-1, i)$, we sum these independent bounds over disjoint temporal windows (e.g., phase pairs $(1, 2), (3, 4), \dots$). This yields a total global cost for OPT over the entire sequence of at least $\sum_j \min(1, \frac{1}{\beta})m_{2j} \geq \frac{1}{2} \min(1, \frac{1}{\beta}) \sum_i m_i = \Omega\left(\min(1, \frac{1}{\beta}) \sum_i m_i\right)$. $\square$

Optimizing the Scaling Parameter β

To determine the optimal competitive ratio $\mathcal{R}$, we balance the upper bound on ALG's cost against OPT's lower bound. Combining Lemma 7 and Lemma 8, and assuming $\beta \geq 1$ and a minimum structural progress of $m_i \geq 1$, we obtain:

$$\mathcal{R} \leq O\left(\beta \min(\log k, \log l) + \frac{k}{\beta}\right)$$

To minimize this ratio, we must perfectly balance the algorithmic tension between structural cache fetches (the $\beta \min(\log k, \log l)$ term) and individual request bypasses (the k/β term). Minimizing this function with respect to β yields the optimal continuous budget scaling parameter:

$$\beta = \sqrt{\frac{k}{\min(\log k, \log l)}}$$

Substituting this optimal β back into our bound perfectly equates the two terms, strictly guaranteeing a final competitive ratio of:

$$\mathcal{R} = O\left(\sqrt{k \min(\log k, \log l)}\right)$$

This result successfully recovers the state-of-the-art upper bound established by [2], while formally generalizing the framework to accommodate fully distinct, individualized request bypass penalties on arbitrary DAG topologies.

6 A Deterministic $O(\min(k, l))$ -Competitive Algorithm for the Weighted Case

We now extend our results to the weighted version of the problem, where each page p has an associated fetching/eviction cost $c(p)$. To achieve an $O(\min(k, l))$ -competitive ratio in the weighted setting, we reduce our problem to an instance of Non-Linear Paging.

6.1 Reduction to Non-Linear Paging

The Non-Linear Paging problem generalizes classic paging by allowing the "size" of a set of pages in the cache to be determined by a custom, monotone function f . A cache state S is considered feasible as long as $f(S) \leq k$.

To reduce our Weighted Dependency Paging problem to Non-Linear Paging, we keep the exact same pages, request sequence, cache capacity k , and eviction costs. The only thing we need to map is the feasibility function f .

For any set of pages S , let its **closure** be the set S combined with all of its required dependencies in the DAG G . We define our feasibility function as the total number of unique pages in that closure:

$$f(S) = |\text{closure}(S)|$$

Because adding more pages to S can only increase or maintain the size of its closure, f is a valid, monotone function (note it is submodular).

Suppose the Non-Linear Paging algorithm outputs cache state S_t to serve the current request. Using our incremental page request assumption, we can assume without loss of generality that S_t satisfies all dependency constraints, since it has no reason to evict a non-maximal page (this action won't change the capacity)

6.2 Analysis of the Width Parameter

The competitive ratio of a non-linear paging algorithm is governed by its width parameter, l_{NL} , defined as the maximum cardinality of a minimally infeasible set minus one.

Lemma 9 (Size of Minimally Infeasible Sets). *Let $\mathcal{M}$ be the collection of all minimally infeasible sets for $f(S) = |\text{closure}(S)|$. For any $S \in \mathcal{M}$, the cardinality $|S|$ is bounded by $\min(k+1, l)$.*

Proof. By definition, a set $S \in \mathcal{M}$ satisfies $f(S) > k$, but any proper subset $S' \subset S$ satisfies $f(S') \leq k$.

First, we show that S must be an antichain in the dependency DAG G . Assume for contradiction that there exist two distinct pages $x, y \in S$ such that x is a dependency of y (i.e., $x \in G[y]$). Because x is already subsumed by the closure of y , removing x from S does not alter the total closure: $\text{closure}(S \setminus \{x\}) = \text{closure}(S)$. This implies $f(S \setminus \{x\}) = f(S) > k$, which directly violates the minimality of S . Thus, no page in S can depend on another page in S , meaning S is an antichain. Since the width of G is l , it must be that $|S| \leq l$.

Second, because the cache capacity is k , and the base cardinality of any individual page's closure is at least 1, the maximum size of a minimally infeasible set cannot exceed $k+1$.

Therefore, the maximum cardinality of $S \in \mathcal{M}$ is strictly bounded by $\min(k+1, l)$. $\square$

Theorem 3. *There is a deterministic $O(\min(k, l))$ -competitive algorithm for the Weighted Dependency Paging problem.*

Proof. By Lemma 9, the maximum cardinality of a minimally infeasible set in our reduction is at most $\min(k+1, l)$. Following the definition of the width parameter in non-linear paging, $l_{NL} = \max_{S \in \mathcal{M}}(|S| - 1)$. This directly implies that $l_{NL} \leq \min(k, l - 1)$.

[3] present a deterministic algorithm that achieves a tight l_{NL} -competitive ratio for general non-linear paging. By deploying their algorithm on our constructed instance, we immediately achieve an $O(l_{NL})$ -competitive ratio, which simplifies to $O(\min(k, l))$. $\square$

7 Deterministic Weighted Dependency Paging with Rejections

In this section, we extend the deterministic weighted dependency paging algorithm to accommodate arbitrary bypass penalties. We build upon the Non-Linear Paging framework, utilizing the closure property of the feasibility function $f(S) = |\text{closure}(S)|$ to inherently manage the dependency DAG topology without requiring fractional routing variables.

7.1 LP Relaxation for Non-Linear Paging with Rejections

We extend the general non-linear LP by introducing the bypass penalty $b_t \in \mathbb{R}^+$ and the bypass variable $z_t \in \{0, 1\}$ for the request at time t . We maintain the definitions of $\mathcal{S}(t) = \{S \subseteq \mathcal{P} \mid p_t \in S \text{ and } f(S) > k\}$, $r(p, t)$, and $I(p, j)$.

Primal LP:

$$\begin{aligned} \min \quad & \sum_{p \in \mathcal{P}} \sum_{j \in [n_p]} x_p(j) \cdot c(p) + \sum_{t \in T} z_t \cdot b_t \\ \text{s.t.} \quad & \sum_{p \in S \setminus \{p_t\}} x_p(r(p, t)) + z_t \geq 1 \quad \forall t \in T, \forall S \in \mathcal{S}(t) \\ & x_p(j) \geq 0, \quad z_t \geq 0 \quad \forall p \in \mathcal{P}, \forall j \in [n_p], \forall t \in T \end{aligned}$$

Dual LP:

$$\begin{aligned} \max \quad & \sum_{t \in T} \sum_{S \in \mathcal{S}(t)} y_t(S) \\ \text{s.t.} \quad & \sum_{t \in I(p, j)} \sum_{S \in \mathcal{S}(t) \mid p \in S \setminus \{p_t\}} y_t(S) \leq c(p) \quad \forall p \in \mathcal{P}, \forall j \in [n_p] \\ & \sum_{S \in \mathcal{S}(t)} y_t(S) \leq b_t \quad \forall t \in T \\ & y_t(S) \geq 0 \quad \forall t \in T, \forall S \in \mathcal{S}(t) \end{aligned}$$

7.2 The Deterministic Algorithm

To achieve the optimal competitive ratio, the algorithm must “revert” tentative evictions if the dual variable for rejection reaches the penalty bound b_t . We denote the accumulated dual weight for page p as $Y_p(j) = \sum_{t \in I(p,j)} \sum_{S \in \mathcal{S}(t) | p \in S \setminus \{p_t\}} y_t(S)$.

Algorithm 4 Deterministic Non-Linear Paging with Rejections

```

1: Initialize primal solutions  $x \leftarrow 0, z \leftarrow 0$  and dual solution  $y \leftarrow 0$ .
2: for time  $t \in T$  do
3:   Let  $Cache_t = \{p \in \mathcal{P} \mid r(p, t) \geq 1 \text{ and } x_p(r(p, t)) = 0\}$ .
4:   Let  $E_t \leftarrow \emptyset$  ▷ Set of tentative evictions at time  $t$ 
5:   while  $f(Cache_t) > k$  do
6:     Find  $Q \subseteq Cache_t$  such that  $Q \in \mathcal{S}(t)$  of minimum cardinality.
7:     while  $Y_p(r(p, t)) < c(p) \forall p \in Q \setminus \{p_t\}$  and  $\sum_{S \in \mathcal{S}(t)} y_t(S) < b_t$  do
8:       Increase  $y_t(Q)$  continuously.
9:     end while
10:    if  $\sum_{S \in \mathcal{S}(t)} y_t(S) = b_t$  then
11:       $z_t \leftarrow 1$  ▷ Request is bypassed
12:      for  $p \in E_t$  do
13:         $x_p(r(p, t)) \leftarrow 0$  ▷ Revert tentative evictions
14:      end for
15:       $Cache_t \leftarrow Cache_t \cup E_t \setminus \{p_t\}$ 
16:      break while loop
17:    end if
18:    for  $p \in Q \setminus \{p_t\}$  such that  $Y_p(r(p, t)) = c(p)$  do
19:      Remove  $p$  from cache:  $x_p(r(p, t)) \leftarrow 1, Cache_t \leftarrow Cache_t \setminus \{p\}$ 
20:       $E_t \leftarrow E_t \cup \{p\}$ 
21:    end for
22:  end while
23: end for
24: Return the solutions  $x, z$ , and  $y$ .

```

Observation 4 (Implicit Topological Enforcement via Closure). *In Algorithm 4, it is possible for a page $p \in Q$ to be evicted while a page x that depends on p remains in $Cache_t$. However, this does not violate physical cache constraints. If $x \in Cache_t$ depends on p , then p is inherently contained within $\text{closure}(Cache_t)$. Therefore, mathematically removing p from $Cache_t$ does not remove p from the closure, meaning the capacity $f(Cache_t)$ does not decrease. Because $f(Cache_t) > k$ remains true, the while loop is forced to continue executing. It will inevitably identify and evict the dependent page x in subsequent iterations to actually reduce the closure size. This mechanism acts as a “lazy eviction”, where the LP absorbs the cost of p immediately, but p ’s physical memory is only freed once its dependents are also cleared.*

7.3 Competitive Analysis

Lemma 10 (Feasibility). *Algorithm 4 returns primal and dual feasible solutions.*

Proof. The dual solution y is strictly bounded by $c(p)$ and b_t due to the halting conditions in Line 7, ensuring dual feasibility. For the primal solution, the while loop at Line 5 only terminates under two conditions:

1. $f(Cache_t) \leq k$, meaning no $S \in \mathcal{S}(t)$ is fully contained in the cache, thus satisfying $\sum_{p \in S \setminus \{p_t\}} x_p(r(p, t)) \geq 1$.
2. $z_t = 1$, which inherently satisfies the constraint $\sum_{p \in S \setminus \{p_t\}} x_p(r(p, t)) + z_t \geq 1$ for all $S \in \mathcal{S}(t)$.

Thus, the primal solution is feasible. □

Lemma 11 (Cost Bound). *Let $c(x, z)$ be the cost of the integral primal solution and $v(y)$ be the value of the dual solution. Then $c(x, z) \leq (l_{NL} + 1) \cdot v(y)$, where l_{NL} is the width of f .*

Proof. Because primal variables are only set to 1 when their corresponding dual constraints are perfectly tight, we have:

$$c(x, z) = \sum_{p \in \mathcal{P}} \sum_{j \in [n_p]} x_p(j) \cdot c(p) + \sum_{t \in T} z_t \cdot b_t \leq \sum_{p \in \mathcal{P}} \sum_{j \in [n_p]} x_p(j) \cdot Y_p(j) + \sum_{t \in T} z_t \sum_{S \in \mathcal{S}(t)} y_t(S)$$

Expanding $Y_p(j)$ and exchanging the order of summation yields exactly:

$$c(x, z) \leq \sum_{t \in T} \sum_{S \in \mathcal{S}(t)} y_t(S) \left(\sum_{p \in S \setminus \{p_t\}} x_p(r(p, t)) + z_t \right)$$

Consider any set S for which $y_t(S) > 0$. By Line 6, S was chosen as a set Q of minimum cardinality, meaning S is minimally infeasible ($S \in \mathcal{M}$). This structural property enforces that $|S \setminus \{p_t\}| \leq l_{NL}$. We evaluate the multiplier in the parentheses:

The variable $x_p(r(p, t))$ evaluates to 1 if page p is evicted at *any* point during the interval $I(p, r(p, t))$.

- **Case 1: The request is served** ($z_t = 0$). The number of terms in the summation is exactly $|S \setminus \{p_t\}| \leq l_{NL}$. Even if every page in the set is eventually evicted, the multiplier is strictly bounded by $l_{NL} + 0 = l_{NL}$.
- **Case 2: The request is bypassed** ($z_t = 1$). At time t , all tentative evictions E_t are reverted. However, pages in $S \setminus \{p_t\}$ may still be legitimately evicted at some future time $t' > t$ before their next request, causing their x_p variables to ultimately become 1. In the absolute worst case, the algorithm pays the bypass penalty ($z_t = 1$) *and* eventually evicts every other page in the minimally infeasible set. The maximum possible sum is mathematically bounded by $1 + |S \setminus \{p_t\}| \leq 1 + l_{NL}$.

In all cases, the multiplier is strictly bounded by $l_{NL} + 1$. Thus:

$$c(x, z) \leq \sum_{t \in T} \sum_{S \in \mathcal{S}(t)} y_t(S) \cdot (l_{NL} + 1) = (l_{NL} + 1) \cdot v(y)$$

□

Theorem 4. *The deterministic primal-dual algorithm for Weighted Dependency Paging with Rejections is $O(\min(k, l))$ -competitive.*

Proof. By Lemma 10, the dual solution y is feasible. By weak duality, $v(y) \leq OPT$. Applying the bound established in Lemma 11:

$$c(x, z) \leq (l_{NL} + 1) \cdot v(y) \leq (l_{NL} + 1) \cdot OPT$$

By Lemma 9, the width l_{NL} of the feasibility function $f(S) = |\text{closure}(S)|$ is strictly bounded by $\min(k, l - 1)$. Therefore, the competitive ratio is bounded by $\min(k + 1, l)$, which simplifies asymptotically to $O(\min(k, l))$. □

8 A Randomized Algorithm for the Weighted Case

In this section, we transition to the weighted version of the randomized dependency paging problem. We first examine a natural fractional relaxation and demonstrate its structural limitations.

8.1 Approach 1: Fractional Relaxation

A standard technique for designing randomized competitive algorithms is to first formulate a continuous, fractional relaxation of the problem state. In our context, we can define the state of the cache by assigning a scalar $y_i \in [0, m]$ to each chain C_i , representing how many of its pages are currently in the cache. We allow these topmost pages to be fractional as long as the total cache capacity is strictly maintained: $\sum_{i=1}^l y_i \leq k$.

Theorem 5. *The fractional relaxation cannot achieve a competitive ratio better than l .*

Proof. Consider l disjoint chains of length m with a cache capacity $k = (m - 1)l$. Let the fetching cost of the h -th page in any chain be $w_h = 2^h$.

Since $\sum_{i=1}^l y_i \leq (m - 1)l$, at any moment there exists at least one chain j such that $y_j \leq m - 1$. The adversary repeatedly requests the m -th page of such a chain ($p_{j,m}$).

To serve $p_{j,m}$, the fractional algorithm must increase y_j to m . Since the page was entirely missing ($y_j \leq m - 1$), the algorithm must fetch $p_{j,m}$ and incur a cost of at least $w_m = 2^m$. To maintain the capacity constraint, the algorithm must decrease the mass of other chains, ensuring the existence of a new chain with $y \leq m - 1$ for the next request. Over a cycle of l requests (one for each chain), the algorithm pays $l \cdot 2^m$.

In contrast, an offline algorithm (OPT) can keep $l - 1$ chains fully in the cache and leave one chain entirely empty. OPT only pays when the empty chain is requested, at which point it swaps it with another chain. For the same l requests, OPT performs at most one swap, costing:

$$\sum_{h=1}^m 2^h < 2 \cdot 2^m$$

The competitive ratio is therefore $\Omega(\frac{l \cdot 2^m}{2 \cdot 2^m}) = \Omega(l)$. $\square$

8.2 Approach 2: Direct Fractional Relaxation with Dependency Constraints

To achieve an optimal $O(\log k)$ competitive ratio independent of general non-linear paging rounding schemes, we formulate a fractional LP that explicitly captures the temporal requests and the topological constraints of the dependency DAG.

Let E be the set of directed edges in the dependency DAG, where $(p, q) \in E$ means page p depends on page q . By the rules of the cache, if dependency q is evicted, its dependent page p must also be evicted. We define our anti-cache variables as $x_{p,j}$, representing the fractional eviction of page p during the interval $I(p, j)$. To enforce dependencies fractionally, we strictly require $x_{p,r(p,t)} \geq x_{q,r(q,t)}$ at any time t .

Let $q(S)$ be the minimum number of pages needed to be excluded from an infeasible set S such that the closure of the remaining cached pages is feasible (i.e., its total size including dependencies is at most k).

8.2.1 The LP Formulation

We define the Primal and Dual LP relaxations explicitly modeling the DAG constraints.

Primal LP:

$$\min \sum_{p \in \mathcal{P}} \sum_{j \in [n_p]} c(p) x_{p,j}$$

Subject to:

$$\sum_{p \in S \setminus \{p_t\}} x_{p,r(p,t)} \geq q(S) \quad \forall t \in T, \forall \text{ infeasible } S \ni p_t$$

$$x_{p,r(p,t)} - x_{q,r(q,t)} \geq 0 \quad \forall (p, q) \in E, \forall t \in T$$

$$x_{p,j} \geq 0 \quad \forall p \in \mathcal{P}, \forall j \in [n_p]$$

Dual LP:

$$\max \sum_{t \in T} \sum_{S \in \mathcal{S}(t)} y_t(S) q(S)$$

Subject to:

$$\sum_{t \in I(p,j)} \left(\sum_{S \in \mathcal{S}(t) | p \in S \setminus \{p_t\}} y_t(S) + \sum_{q: (p,q) \in E} \gamma_{p,q,t} - \sum_{r: (r,p) \in E} \gamma_{r,p,t} \right) \leq c(p) \quad \forall p \in \mathcal{P}, \forall j \in [n_p]$$

$$y_t(S) \geq 0, \gamma_{p,q,t} \geq 0$$

The dual variables $\gamma_{p,q,t}$ act as a directed "flow" along the DAG. If the algorithm is forced to fractionally evict a dependency q , it can shift the dual cost burden upwards to the dependent page p that requires it, preventing q 's dual constraint from being prematurely violated.

8.2.2 The Fractional Algorithm

We maintain a fractional state x and actively group pages into disjoint **components** K . We now explicitly manage the γ variables within these components to route dual flow and maintain strict dual feasibility while enforcing $x_p \geq x_q$.

1. **Initialization:** Start with $x_{p,j} = 0$ for all pages. Initialize all dual variables $y_t(S) = 0$ and internal flow variables $\gamma_{p,q,t} = 0$. Let the set of components $\mathcal{K}$ initially be individual pages: $\mathcal{K} = \{\{p\} : p \in \mathcal{P}\}$.
2. **On Request:** When a request p_t arrives at time t , $x_{p_t,r(p_t,t)} = 0$. Any existing component containing p_t or its dependencies is broken down back into individual pages, and their associated internal γ flows are frozen for that interval.
3. **Continuous Update:** While there is an infeasible set $S \ni p_t$ such that the primal constraint $\sum_{p \in S \setminus \{p_t\}} x_{p,r(p,t)} < q(S)$ is violated, and $x_{p,r(p,t)} < 1$ for all $p \in S \setminus \{p_t\}$:
 - Continuously increase the dual variable $y_t(S)$.
 - For each page p , its native accumulated dual weight is $Y_{p,r(p,t)} = \sum_{\tau \leq t} \sum_{S' \ni p} y_\tau(S')$.
 - For each component $K \in \mathcal{K}$, its aggregated weight is $Y_K = \sum_{p \in K} Y_{p,r(p,t)}$.
 - **Dual Flow Routing (γ update):** Within any merged component K , the dual weight must be distributed so no page violates its capacity. We continuously increase the internal flow variables $\gamma_{p,q,t}$ along the directed edges of K . Specifically, for an infinitesimal increase dY_K , we route γ such that the *effective* dual weight increment dY_v^{eff} for every page $v \in K$ satisfies:

$$dY_v^{\text{eff}} = dy_t(S) \cdot \mathbb{I}[v \in S] + \sum_{(v,u) \in E} d\gamma_{v,u,t} - \sum_{(r,v) \in E} d\gamma_{r,v,t} = \frac{c(v)}{c(K)} dY_K$$

where $c(K) = \sum_{p \in K} c(p)$. This explicitly shifts excess dual weight from cheaper dependencies to expensive dependents.

- Increase the uniform anti-cache value x_K (applied identically to all $p \in K$) exponentially using the component's total capacity:

$$x_K = \frac{1}{k} \left(\exp \left(\frac{\ln(k+1)}{c(K)} Y_K \right) - 1 \right)$$

- **Component Merging:** As dual variables grow, if x_q reaches x_p for some dependency edge $(p,q) \in E$ and x_q is growing faster, we merge them to prevent a topological violation. We replace K_p and K_q with a new component $K' = K_p \cup K_q$. Their fractional values are locked together, permanently enforcing $x_{p,r(p,t)} \geq x_{q,r(q,t)}$ and allowing γ flow between them.

Observation 5 (Continuity of Component Merging). *When two components K_p and K_q merge, their fractional value x does not jump. The merge occurs exactly when their dual densities are equal: $\frac{Y_{K_p}}{c(K_p)} = \frac{Y_{K_q}}{c(K_q)} = D$. By the mediant property of fractions, the combined density is $\frac{Y_{K_p} + Y_{K_q}}{c(K_p) + c(K_q)} = D$, strictly preserving the exponent and ensuring continuous evolution.*

8.2.3 Competitive Ratio Analysis

Theorem 6. *The fractional algorithm yields a fractional $O(\log k)$ -competitive ratio for the Weighted Dependency Paging problem.*

Proof. Step 1: Dual Feasibility and Network Flow. To utilize weak duality, we must prove the constructed dual variables y and γ form a feasible solution to the Dual LP. The left-hand side of the dual constraint for any page v is exactly its effective dual weight Y_v^{eff} .

For an individual unmerged page p , $\gamma = 0$, and the stopping condition $x_p \leq 1$ trivially ensures $\ln(k+1) \geq \frac{\ln(k+1)}{c(p)} Y_p^{\text{eff}}$, meaning $Y_p^{\text{eff}} \leq c(p)$.

For a merged component K , the algorithm explicitly routes the γ variables such that $Y_v^{\text{eff}} = \frac{c(v)}{c(K)} Y_K$. Because the algorithm guarantees $x_K \leq 1$ for the entire component, substituting this into the update rule gives:

$$1 \geq \frac{1}{k} \left(\exp \left(\frac{\ln(k+1)}{c(K)} Y_K \right) - 1 \right) \implies Y_K \leq c(K)$$

Since $Y_K \leq c(K)$, the effective dual weight for each individual page $v \in K$ is strictly bounded:

$$Y_v^{\text{eff}} = \frac{c(v)}{c(K)} Y_K \leq \frac{c(v)}{c(K)} c(K) = c(v)$$

Because dependencies q merge with dependents p only when q 's density exceeds p 's, the required flow direction inherently aligns with the valid edge orientation (p, q) (meaning q successfully offloads weight $-\gamma$ while p absorbs $+\gamma$). Thus, the explicit γ routing guarantees every dual constraint is strictly satisfied.

Step 2: Bounding the Derivative. Let P denote the fractional primal cost and D denote the dual cost. An infinitesimal increase $dy_t(S)$ increases the dual cost by $dD = q(S)dy_t(S)$. Taking the derivative of our exponential update rule with respect to Y_K gives $c(K)dx_K = \ln(k+1)(x_K + 1/k)dY_K$.

Since $dY_K = |K \cap S|dy_t(S)$, the total increase in our primal cost is:

$$dP = \sum_{K \in \mathcal{K}} c(K)dx_K = \ln(k+1)dy_t(S) \sum_{K \in \mathcal{K}} |K \cap S| \left(x_K + \frac{1}{k} \right)$$

Summing $|K \cap S|$ exactly reconstructs the individual pages in the set, allowing the grouped component costs $c(K)$ to drop out:

$$dP = \ln(k+1)dy_t(S) \sum_{p \in S \setminus \{p_t\}} \left(x_{p,r(p,t)} + \frac{1}{k} \right)$$

Step 3: Utilizing the Cache Capacity. By definition, excluding the optimal $q(S)$ pages from S results in a set S^* such that its closure is feasible ($|\text{closure}(S^*)| \leq k$). Since $S^* \subseteq \text{closure}(S^*)$, the raw remaining pages satisfy $|S| - q(S) \leq k$, or $\frac{|S|}{k} \leq \frac{q(S)}{k} + 1 \leq 2q(S)$.

During the update, the primal constraint is unmet ($\sum_{p \in S \setminus \{p_t\}} x_{p,r(p,t)} \leq q(S)$). Substituting this structural bound into the derivative:

$$dP \leq \ln(k+1)dy_t(S) \left(\sum_{p \in S \setminus \{p_t\}} x_{p,r(p,t)} + \frac{|S| - 1}{k} \right)$$

$$dP \leq \ln(k+1)dy_t(S) (q(S) + 2q(S)) = 3\ln(k+1) \cdot q(S)dy_t(S)$$

Conclusion. Substituting $dD = q(S)dy_t(S)$, we establish $dP \leq 3\ln(k+1) \cdot dD$. Integrating this over the request sequence yields a total fractional primal cost strictly bounded by the final dual cost:

$$P \leq 3\ln(k+1) \cdot D$$

Because Step 1 established explicit dual feasibility via the γ routing, weak duality perfectly applies: $D \leq \text{OPT}$. Therefore, the fractional cost is bounded by $P \leq 3\ln(k+1) \cdot \text{OPT}$, yielding an expected $O(\log k)$ fractional competitive ratio. $\square$

8.2.4 The Necessity of the Subset Constraints

One might wonder why the Primal LP requires an exponential number of constraints—one for every infeasible subset $S \in \mathcal{S}(t)$ —rather than simply enforcing a single global cache capacity constraint alongside the explicit dependency constraints ($x_p \geq x_q$).

The Global Capacity Constraint: Notice that our formulation inherently includes the global capacity constraint. If we let $S = \mathcal{P}$ (the set of all pages in the system), the number of pages we must exclude to reach the cache capacity k is $q(\mathcal{P}) = |\mathcal{P}| - k$. The subset constraint for $S = \mathcal{P}$ becomes $\sum_{p \in \mathcal{P}} x_p \geq |\mathcal{P}| - k$, which seamlessly rearranges to the standard capacity limit: $\sum_{p \in \mathcal{P}} (1 - x_p) \leq k$.

However, relying strictly on the global capacity constraint introduces severe structural violations and massive integrality gaps, as demonstrated by the following two counterexamples.

Counterexample 1: Hiding Fractional Mass

Consider a cache of capacity k . There is a set B of $k-1$ “bottom” pages, and a set T of $k-1$ “top” pages. The dependencies are configured such that *every* top page $t \in T$ depends on *all* $k-1$ bottom pages in B .

In any valid integral solution, caching even a single top page t requires caching all $k-1$ bottom pages. This consumes $1 + (k-1) = k$ slots, perfectly filling the cache. Therefore, **every integer solution must include at most 1 top page**, meaning at least $k-2$ top pages must be entirely evicted ($x_t = 1$).

Suppose we formulate a simplified LP using *only* the global capacity constraint ($\sum_{p \in \mathcal{P}} x_p \geq (2k-2) - k = k-2$) and the dependency constraints ($x_t \geq x_b$ for all $t \in T, b \in B$).

A fractional solution could simply assign $x_p = 0.5$ to all pages in the system.

- **Global capacity is satisfied:** $\sum_{p \in \mathcal{P}} x_p = 2(k-1) \times 0.5 = k-1$. Since $k-1 \geq k-2$, the constraint holds.
- **Dependencies are satisfied:** $x_t \geq x_b \implies 0.5 \geq 0.5$.

However, this solution leaves $1 - 0.5 = 0.5$ fractional mass of *every* top page inside the cache. The total fractional mass of cached top pages is $\frac{k-1}{2}$. For $k > 3$, this vastly exceeds the logical maximum of 1, allowing the algorithm to “cheat” by hiding an impossible combination of fractional top pages inside the global capacity pool.

Counterexample 2: Cost Smearing and the $\Omega(k)$ Integrality Gap

Now consider the weighted setting. Let the cache capacity be $k = m$. Consider a strict dependency chain C consisting of m pages: $p_1 \leftarrow p_2 \leftarrow \dots \leftarrow p_m$, where each p_i depends on p_{i-1} .

- The fetching cost for the lower pages is $c(p_i) = 1$ for all $1 \leq i \leq m-1$.
- The fetching cost for the topmost page is extremely large: $c(p_m) = W$ (where $W \gg m$).
- There is one additional independent page L with a fetching cost of $c(L) = 1$.

The total number of pages is $m+1$. Assume C is fully cached. The adversary repeatedly requests L , then p_m .

For an **integral algorithm** to serve L , it must evict a page to respect capacity m . Evicting any page from C forces the eviction of p_m due to dependency rules. When p_m is requested next, the algorithm pays W to refetch it. Thus, the integral optimal pays $\Omega(W)$ per phase.

Under the **simplified fractional LP** (global constraint $\sum_{p \in \mathcal{P}} x_p \geq 1$ and dependencies $x_{p_m} \geq \dots \geq x_{p_1}$), when L is requested, the algorithm must set $x_L = 0$. To satisfy the global eviction requirement of 1 as cheaply as possible, it can “smear” the eviction uniformly across the entire chain by setting $x_{p_i} = \frac{1}{m}$ for all $1 \leq i \leq m$. This satisfies the global capacity ($\sum_{i=1}^m \frac{1}{m} = 1$) and the dependencies ($\frac{1}{m} \geq \frac{1}{m}$). When p_m is requested next, the fractional algorithm pays to refetch this fraction:

$$\text{Fractional Cost} = \sum_{i=1}^m c(p_i) x_{p_i} = \frac{1}{m} \left((m-1) \cdot 1 + W \right) \approx \frac{W}{m}$$

By exploiting the global constraint, the fractional algorithm pays roughly W/m per phase compared to the integral W . This results in an integrality gap of $\Omega(m)$. Since $m = k$, this $\Omega(k)$ gap completely breaks any expected $O(\log k)$ competitive ratio.

How Our Subset Constraints Fix the Issues

Our full LP enforces $\sum_{p \in S} x_p \geq q(S)$ for *every* infeasible subset, addressing both exploits simultaneously.

Fixing Counterexample 1: Consider the subset $S = T$ (the $k-1$ top pages). To make any subset of T feasible, its closure must be at most k . Because caching even one top page brings in all $k-1$ bottom pages, we can retain at most 1 top page. To make $S = T$ feasible, we must exclude $q(T) = |T| - 1 = k-2$ pages. This generates the constraint:

$$\sum_{t \in T} x_t \geq k-2$$

If we evaluate the cheating fractional solution ($x_t = 0.5$), the sum is $\frac{k-1}{2}$. For $k = 5$, this evaluates to 2, which strictly violates the required bound of $k-2 = 3$.

Fixing Counterexample 2: Consider the subset $S = \{p_m, L\}$. Caching both brings the entire chain C into the cache, giving a closure size of $m+1 > k$. Thus, S is infeasible. Excluding either p_m or L leaves a feasible closure ($\leq m$). Therefore, $q(S) = 1$. This generates the constraint:

$$x_{p_m} + x_L \geq 1$$

When the adversary requests L , the algorithm is forced to set $x_L = 0$. This subset constraint immediately enforces:

$$x_{p_m} \geq 1$$

The fractional algorithm is mathematically blocked from smearing its eviction mass into the cheaper lower pages. It is forced to fully evict the expensive topmost page p_m , paying the full cost W to refetch it and completely closing the integrality gap.

9 Fractional Relaxation via Knapsack Cover Inequalities

To achieve an optimal fractional approximation for the weighted dependency paging problem, we formulate a relaxation based on Knapsack Cover (KC) inequalities adapted to the poset structure of the dependencies.

Let v be a vertex in the dependency graph, c_v be the intrinsic fetching/eviction cost of that vertex, and K be the cache capacity. We assume dependencies are structured such that to cache a leaf, we must cache its ancestors (dependencies flow up the tree).

Definition (Subtree Weight): The weight of a vertex, W_v , is the sum of the costs in its subtree (the vertex and all its descendants/dependencies):

$$W_v = \sum_{u \in \text{subtree}(v)} c_u$$

Definition (Antichain Weight): For any antichain A (a set of mutually incomparable vertices in the dependency poset), the weight of the antichain, W_A , is the sum of the subtree weights of its vertices:

$$W_A = \sum_{v \in A} W_v$$

Definition (Adjusted Weight): For any vertex $v \in A$, the adjusted weight $\widetilde{W}_v(A)$ is defined to cap the weight contribution at the required capacity to clear:

$$\widetilde{W}_v(A) = \min(W_v, W_A - K)$$

The Fractional Constraint: Let $x_v \in [0, 1]$ be the fractional eviction variable for vertex v . For every antichain A where $W_A > K$, the fractional solution must clear enough space to satisfy the capacity:

$$\sum_{v \in A} \widetilde{W}_v(A) \cdot x_v \geq W_A - K$$

9.1 Redundancy of Explicit Dependency Constraints

A standard fractional formulation would require explicit constraints to enforce the tree structure: if ancestor v depends on leaf u , then evicting u forces the eviction of v (i.e., $x_v \geq x_u$). However, we can prove that our modified KC inequalities structurally enforce these dependencies naturally, rendering explicit constraints mathematically redundant without loss of generality.

Theorem 7 (Implicit Dependency Enforcement). *The fractional LP defined by the modified KC inequalities over all antichains is naturally monotonically constrained such that an optimal solution organically satisfies $x_v \geq x_u$ for any ancestor v and its descendant u .*

Proof. Let u be a descendant and v be its ancestor. Assume the objective is to minimize fractional eviction cost. We prove that for any antichain containing u , there is a corresponding antichain containing v that is strictly harder to cover, forcing the LP to assign $x_v \geq x_u$.

Step 1: Constructing the Corresponding Antichain. Let A_u be any antichain containing u . Let $D \subseteq A_u$ be the set of all elements in A_u that are descendants of v (this includes u). We construct a new set: $A_v = (A_u \setminus D) \cup \{v\}$. Because A_u is an antichain, it cannot contain any ancestors of v . Therefore, swapping the descendants for the ancestor creates a valid antichain A_v .

Step 2: The RHS Inequality. By definition, the weight of an ancestor includes the weight of its subtree. Therefore, $W_v \geq \sum_{d \in D} W_d$. Calculating the total weight of the antichains:

$$W_{A_v} = W_{A_u} - \sum_{d \in D} W_d + W_v \implies W_{A_v} \geq W_{A_u}$$

Let $R_u = W_{A_u} - K$ and $R_v = W_{A_v} - K$ be the space that needs to be cleared (the Right-Hand Side of the KC inequalities). Because the ancestor's antichain is heavier, it requires more space to be cleared: $R_v \geq R_u$.

Step 3: The Relative Covering Efficiency. We write the KC inequalities for both antichains and normalize them by dividing by their respective RHS.

For A_u :

$$\sum_{y \notin D} \frac{\min(W_y, R_u)}{R_u} x_y + \dots + \frac{\min(W_u, R_u)}{R_u} x_u \geq 1$$

For A_v :

$$\sum_{y \notin D} \frac{\min(W_y, R_v)}{R_v} x_y + \frac{\min(W_v, R_v)}{R_v} x_v \geq 1$$

Notice the coefficient of x_v in the second equation. Because $W_v \geq W_{A_v} - K$, we know $W_v \geq R_v$. Therefore, $\min(W_v, R_v) = R_v$, and the coefficient of x_v is exactly 1.

Now look at the external independent variables x_y . Their relative covering efficiency is defined by the function $f(R) = \frac{\min(W_y, R)}{R}$.

- If $R \leq W_y$, $f(R) = 1$.
- If $R > W_y$, $f(R) = \frac{W_y}{R}$, which strictly decreases as R increases.

Because $R_v \geq R_u$, the efficiency of any external variable y is monotonically non-increasing: $f(R_v) \leq f(R_u)$.

Step 4: The Conclusion. Because external variables x_y are algebraically *less effective* at covering the ancestor's antichain A_v than the descendant's antichain A_u , the LP cannot rely on them to clear the required space. The LP must make up the deficit using x_v . This structural asymmetry guarantees two things:

1. **Natural Monotonicity:** To minimize the overall cost, the LP is algebraically forced to push x_v at least as high as x_u , organically satisfying $x_v \geq x_u$.
2. **Protection against Online Requests:** If the environment requests the ancestor ($x_v = 0$), the LP must satisfy the A_v constraint entirely using external variables x_y . Because $R_v \geq R_u$, those external variables will be forced so high that they provide a massive *surplus* of cover for the easier A_u constraint, allowing the LP to naturally drop x_u to 0 without violating the cache capacity limits.

□

References

- [1] Marcin Bienkowski et al. “Online Tree Caching”. In: *arXiv* (2016). DOI: [10.48550/arxiv.1602.08563](https://doi.org/10.48550/arxiv.1602.08563).
- [2] Julien Dallot et al. “Dependency-aware online caching”. In: *IEEE INFOCOM 2024-IEEE Conference on Computer Communications*. IEEE, 2024, pp. 871–880.
- [3] Ilan Doron-Arad, Joseph, and Naor. “Non-Linear Paging”. In: *arXiv* (2024). DOI: [10.48550/arxiv.2404.13334](https://doi.org/10.48550/arxiv.2404.13334).
- [4] Leah Epstein et al. “Online File Caching with Rejection Penalties”. In: *Algorithmica* 71 (2013), pp. 279–306. DOI: [10.1007/s00453-013-9793-0](https://doi.org/10.1007/s00453-013-9793-0).